

80 minutes

Advice: Skip problems that might take a while and come back to them later. If an algorithm is requested, be as succinct as possible. Often a simple description in English is more effective than pseudocode. You do not need to supply any proofs unless explicitly asked. All runtimes should be given using big-oh notation.

1. Simplify the following expressions if possible. [Points: _____/12]

- (a) $O(n^2 + 500n)$.
- (b) $O(n \log n + n^{1.01})$.
- (c) $O(3n^3 - 2n^2 + n)$.
- (d) $O(\sqrt{n} + \log n)$.
- (e) $O(1 + 1/n)$.
- (f) $O(\log(n^2))$.

2. The Mergesort algorithm described in class works as follows on an unordered list U of n numbers:

- If $n = 1$, return U (it is already sorted).
- Break U into two lists U_1 and U_2 (of roughly the same length).
- Recursively sort: $S_1 = \text{Mergesort}(U_1)$ and $S_2 = \text{Mergesort}(U_2)$.
- Output merger: $S = \text{Merge}(S_1, S_2)$

Answer the following questions: [Points: _____/2]

- (a) Give the recurrence relationship that describes the runtime of Mergesort.
 - (b) Give the runtime of Mergesort using big-oh notation.
3. Define 3-Mergesort to break the list into three sublists of length roughly $n/3$, recursively sort these sublists, and then 3-Merge them together. [Points: _____/8]
- (a) Give an algorithm for 3-Merge.
 - (b) Give the runtime for 3-Merge using big-oh notation.
 - (c) Give the recurrence relationship that describes the runtime of 3-Mergesort.
 - (d) Give the runtime for 3-Mergesort using big-oh notation.
4. Define k -Mergesort to break the list into k sublists of length roughly n/k , recursively sort these sublists, and then k -Merge them together. [Points: _____/6]
- (a) Assume that k -Merge (on lists of size n/k) can be implemented in $\Theta(n \log k)$ and give the recurrence relationship for the runtime of k -Mergesort.
 - (b) Give the runtime for k -Mergesort using big-oh notation.
 - (c) Give an algorithm for k -Merge that runs in $O(n \log k)$ time. (Hint: use priority queues.)

5. You are given a graph $G(V, E)$ with edge costs $c(e)$ for $e \in E$. You may assume that the edge costs $c(\cdot)$ are distinct. Let $T \subset E$ be the minimum spanning tree. Consider what happens to the tree if we change the cost of an edge e' . Formally, the new edge costs are $c'(\cdot)$ given by $c'(e') \neq c(e')$ and $c'(e) = c(e)$ if $e \neq e'$. You may assume that the edge costs $c'(\cdot)$ are each distinct. Let T' be the minimum spanning tree with respect to costs $c'(\cdot)$. [Points: _____/10]
- Consider decreasing the cost of an edge in T (i.e., $e' \in T$ and $c'(e') < c(e')$). Is $T' = T$ always?
 - Consider increasing the cost of an edge in T (i.e., $e' \in T$ and $c'(e') > c(e')$). Is $T' = T$ always?
 - Consider decreasing the cost of an edge not in T (i.e., $e' \notin T$ and $c'(e') < c(e')$). Is $T' = T$ always? If not, give an illustrative example that shows why and give a simple algorithm for computing T' from T without solving the MST problem over from scratch.
 - Consider increasing the cost of an edge not in T (i.e., $e' \notin T$ and $c'(e') > c(e')$). Is $T' = T$ always? If not, give an illustrative example that shows why and give a simple algorithm for computing T' from T without solving the MST problem over from scratch.
6. You've just started consulting for a startup company, DigiDyne, that is doing dynamic pricing of digital music downloads. They are considering two business models. In the *subscription* model customers are asked to pay a fixed price p and can download as many songs as they please. In the *a-la-carte* model a customer is asked to pay a fixed price q per song. The price for d downloads in the a-la-carte model is $d \cdot q$.

DigiDyne has done some market research that suggests that each consumer behaves in the following way. Consumer i wishes to download d_i songs and pay at most v_i for the privilege. If the total price consumer i is asked to pay for d_i downloads is at most v_i , they will pay the asked price. If the total price consumer i is asked to pay is more than v_i , they will not pay for any service (Perhaps they will use a competing service instead). Thus, the input to DigiDyne's pricing problem is completely specified by two n -dimensional vectors, $\mathbf{v} = (v_1, \dots, v_n)$ and $\mathbf{d} = (d_1, \dots, d_n)$. [Points: _____/8]

Example:

- $S = \{1, 2, 3\}$, $\mathbf{v} = (4, 5, 6)$, $\mathbf{d} = (1, 2, 3)$.
 - For subscription price $p = 5$: consumers 2 and 3 buy (v_2 and v_3 are greater than $p = 5$), consumer 1 does not buy ($v_1 < p = 5$). The total revenue is 10.
 - For a-la-carte price $q = 3$: consumer 1 buys ($v_1/d_1 \geq q = 3$), consumers 2 and 3 do not buy (v_2/d_2 and $v_3/d_3 < q = 3$). The total revenue is 3 (consumer 1 buys one song for $q = 3$).
- Show that neither of these business models is always better than the other. To do so, give an input (\mathbf{v}, \mathbf{d}) where the revenue from the optimal subscription price, p^* , is less than the revenue from the optimal a-la-carte price, q^* . Then given an input $(\mathbf{v}', \mathbf{d}')$ where the revenue from the optimal subscription price, p^* , is more than the revenue from the optimal a-la-carte price, q^* .
 - Give an algorithm that on input (\mathbf{v}, \mathbf{d}) computes the a-la-carte price, q^* , with the highest total revenue. What is the runtime of your algorithm?